

Case Study 1 - Beers and Breweries in USA

Yucheol Shin, Tai Chowdhury, Patricia Attah

6/21/2020

```
#install.packages("tinytex")  
#tinytex::install_tinytex()
```

Prepare Data

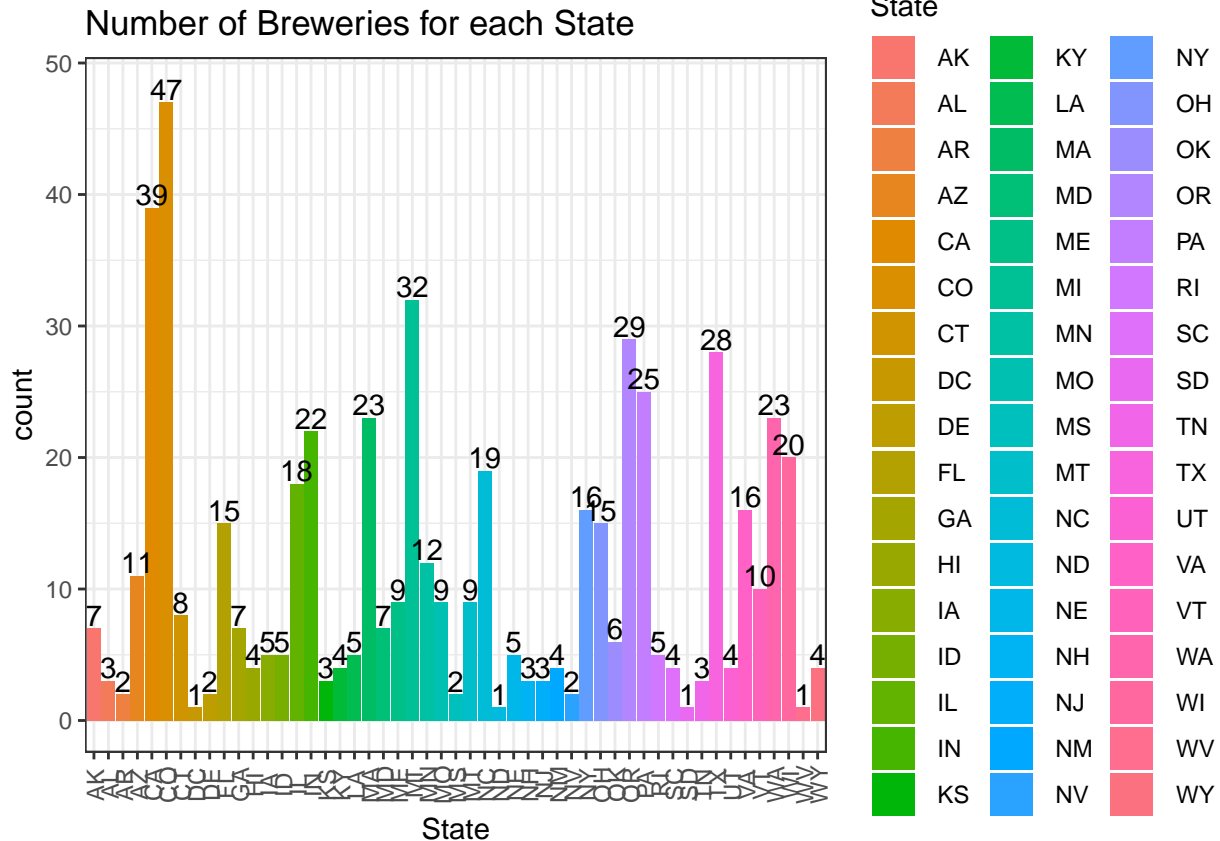
We first prepare data. We import Brew and Breweries data from CSV files and include necessary libraries for our code.

```
library(ggplot2)  
library(tidyr)  
library(plyr)  
library(dplyr)  
library(class)  
library(caret)  
library(e1071)  
library(RCurl)  
library(httr)  
library("RColorBrewer")  
  
x<-getURL('https://raw.githubusercontent.com/yuchrishin/DS6306-GroupProject1/master/data/Beers.csv')  
beers =read.csv(text=x)  
  
y<- getURL('https://raw.githubusercontent.com/yuchrishin/DS6306-GroupProject1/master/data/Breweries.csv')  
breweries = read.csv(text=y)
```

Number of Breweries per state

We want to analyze number of breweries in each states. We count the number of breweries for each state and plot it on top of each bar. From the graph, we can see that Colorado and California have most number of breweries. We see that some states have only 1 breweries such as DC, North and South Dakota. From this graph, we can ask a question why some states have more breweries than others.

```
totalState = count(breweries, State)  
breweries %>% arrange() %>% ggplot(aes(x=State, fill = State)) + geom_bar() + geom_text(aes(State, n + 1))
```



Merge two data

Breweries and Beer data are two separate data. Merging these two datas will give more variables to analyze. For example, we can look into the relationship between states and beers. In order to merge, we need to find if they have key variable that we can join together. Breweries data has Brew_ID and Beer data has Brewery_id which we can merge. Converting the name of column in Beer, two datas are merged as below.

```
colnames(beers)[5] = "Brew_ID"
fullData = merge(beers, breweries, by = "Brew_ID")
head(fullData)
```

##	Brew_ID	Name.x	Beer_ID	ABV	IBU	Style	Ounces	Name.y
## 1	1	Get Together	2692	0.045	50	American IPA	16	NorthGate Brewing
## 2	1	Maggie's Leap	2691	0.049	26	Milk / Sweet Stout	16	NorthGate Brewing
## 3	1	Wall's End	2690	0.048	19	English Brown Ale	16	NorthGate Brewing
## 4	1	Pumpkin	2689	0.060	38	Pumpkin Ale	16	NorthGate Brewing
## 5	1	Stronghold	2688	0.060	25	American Porter	16	NorthGate Brewing
## 6	1	Parapet ESB	2687	0.056	47	Extra Special / Strong Bitter (ESB)	16	NorthGate Brewing

Missing values

In order to process the analysis, we need to clean up data as there might be some missing data or incorrectly formatted data. Below are the code that we have ran to find out if there is any missing data.

```
sapply(fullData, function(x) sum(is.na(x)))
```

```
## Brew_ID Name.x Beer_ID ABV IBU Style Ounces Name.y City State
##      0      0      0   62 1005      0      0      0      0      0
```

```
cleanData = fullData %>% filter(!is.na(ABV) & !is.na(IBU))
```

There are 1005 rows of data that do not have IBU value. We need a IBU data in order to make an analysis so we decided to drop the rows that are missing IBU and ABV data.

Median ABV and IBU per states

We want to look at the median values for each state. In order to get median for each states, we collecte the data grouping by state and summarize them. With the data we calculated, we draw a bar charts with median ABV and IBU of all states.

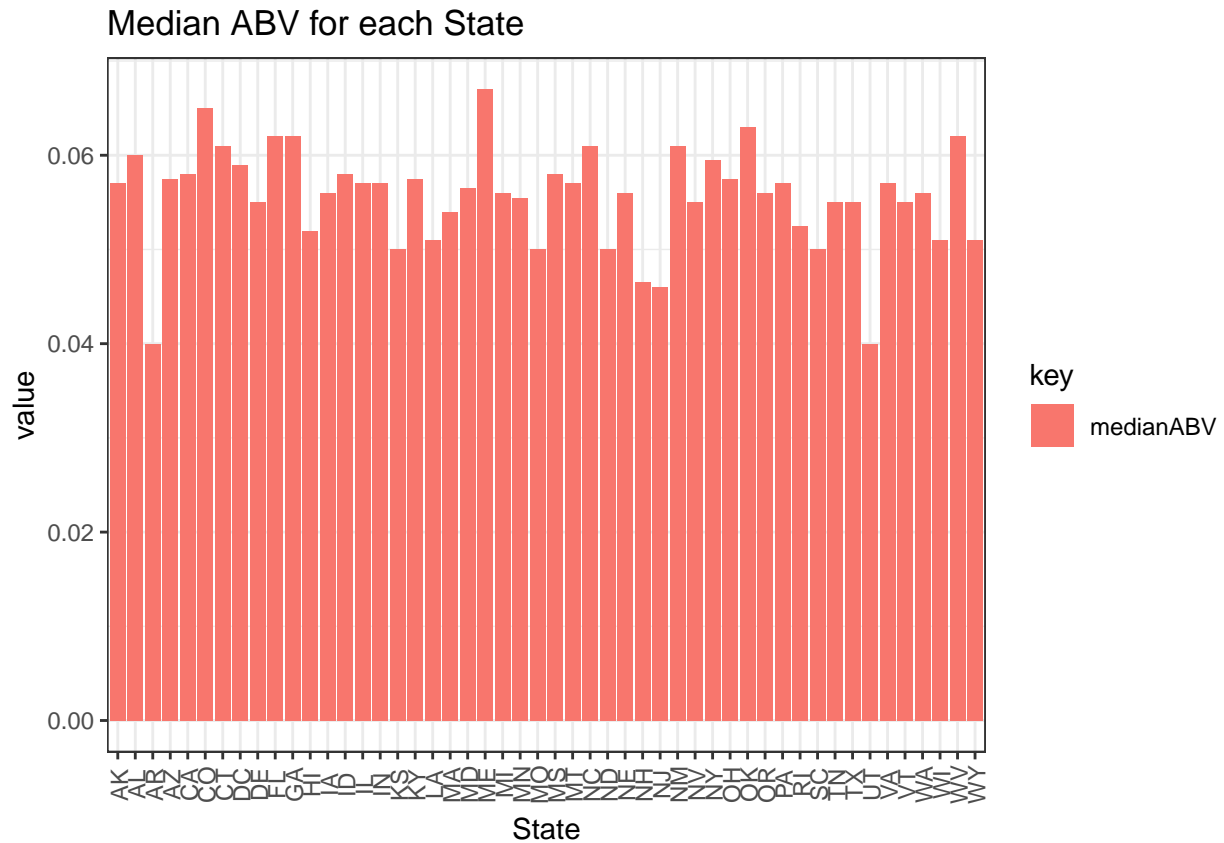
```
cleanData %>% group_by(State) %>% summarize(medianABV = median(ABV), medianIBU = median(IBU), count = n
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 50 x 4
##   State medianABV medianIBU count
##   <fct>      <dbl>      <dbl> <int>
## 1 " AK"      0.057         46      17
## 2 " AL"      0.06         43       9
## 3 " AR"      0.04         39       1
## 4 " AZ"      0.0575        20.5      24
## 5 " CA"      0.058         42     135
## 6 " CO"      0.065         40     146
## 7 " CT"      0.061         29       6
## 8 " DC"      0.059         47.5       4
## 9 " DE"      0.055         52       1
## 10 " FL"     0.062         55      37
## # ... with 40 more rows
```

```
cleanData %>%
  group_by(State) %>%
  summarise(medianABV = median(ABV)) %>%
  gather(key, value, -State) %>%
  ggplot(aes(State, value, fill = key)) + geom_bar(stat = "identity", position = "dodge") + ggtitle("Me
```


```
## 'summarise()' ungrouping output (override with '.groups' argument)
```



```
cleanData %>%
  group_by(State) %>%
  summarise(medianIBU = median(IBU)) %>%
  gather(key, value, -State) %>%
  ggplot(aes(State, value, fill = key)) + geom_bar(stat = "identity", position = "dodge") + ggtitle("Median IBU for each State")

## 'summarise()' ungrouping output (override with '.groups' argument)
```

Median IBU for each State



```
medABV <- cleanData %>% group_by(State) %>% summarize(medianABV = median(ABV), medianIBU = median(IBU),
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
medABV <- medABV %>% select(State, medianABV)
```

```
medIBU <- cleanData %>% group_by(State) %>% summarize(medianABV = median(ABV), medianIBU = median(IBU),
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
medIBU <- medIBU %>% select(State, medianIBU)
```

```
head(medABV)
```

```
## # A tibble: 6 x 2
##   State medianABV
##   <fct>      <dbl>
## 1 " ME"      0.067
## 2 " CO"      0.065
## 3 " OK"      0.063
## 4 " FL"      0.062
## 5 " GA"      0.062
## 6 " WV"      0.062
```

```
tail(medABV)
```

```
## # A tibble: 6 x 2
##   State medianABV
##   <fct>      <dbl>
## 1 " ND"      0.05
## 2 " SC"      0.05
## 3 " NH"     0.0465
## 4 " NJ"     0.046
## 5 " AR"     0.04
## 6 " UT"     0.04
```

```
head(medIBU)
```

```
## # A tibble: 6 x 2
##   State medianIBU
##   <fct>      <dbl>
## 1 " ME"       61
## 2 " WV"     57.5
## 3 " FL"      55
## 4 " GA"      55
## 5 " DE"      52
## 6 " NM"      51
```

```
tail(medIBU)
```

```
## # A tibble: 6 x 2
##   State medianIBU
##   <fct>      <dbl>
## 1 " RI"       24
## 2 " HI"     22.5
## 3 " WY"      21
## 4 " AZ"     20.5
## 5 " KS"      20
## 6 " WI"      19
```

In median ABV bar chart, we can see it quite evenly spread out except for Arizona and Utah that it has significantly lower medians than others. We see that Maine and Colorado has much higher median ABV than other states. For median IBU bar chart, results come out to be distributed wider than median ABV. We see there is dramatic differences for each states. We found Maine and West Virginia have highest median IBU, and Kansas and Wisconsin have lowest median IBU.

Max ABV and IBU of state

In order to get the max ABV and IBU value, we followed two approaches. One is to get max values for each state by grouping state and summarizing each state. Another approach is to get max values among all states.

```
cleanData %>% group_by(State) %>% summarize(maxABV = max(ABV), maxIBU = max(IBU))
```

```
## 'summarise()' ungrouping output (override with '.groups' argument)
```

```
## # A tibble: 50 x 3
##   State maxABV maxIBU
##   <fct> <dbl> <int>
## 1 " AK"  0.065     71
## 2 " AL"  0.093    103
## 3 " AR"  0.04      39
## 4 " AZ"  0.095     99
## 5 " CA"  0.099    115
## 6 " CO"  0.099    104
## 7 " CT"  0.088     85
## 8 " DC"  0.092    115
## 9 " DE"  0.055     52
## 10 " FL" 0.082     82
## # ... with 40 more rows
```

```
maxABV = max(cleanData$ABV)
maxIBU = max(cleanData$IBU)
```

```
cleanData %>% filter(ABV == maxABV)
```

```
##   Brew_ID      Name.x Beer_ID  ABV IBU      Style Ounces      Name.y
## 1         2 London Balling   2685 0.125  80 English Barleywine    16 Against the Grain Brewery Louis
```

```
cleanData %>% filter(IBU == maxIBU)
```

```
##   Brew_ID      Name.x Beer_ID  ABV IBU      Style Ounces      Name.y
## 1        375 Bitter Bitch Imperial IPA    980 0.082 138 American Double / Imperial IPA    12 Astoria I
```

First chart display max ABV and IBU for each state. From the data, we found that London Balling has ABV of 0.125 and it has maximum ABV among all beers. We also found Bitter Bitch Imperial IPA contains IBU of 138 which is the maximum among all beers.

Summarize ABV

Checking distribution of the data is one of key part of EDA. We plot several distributions graphs of ABV to check its normality.

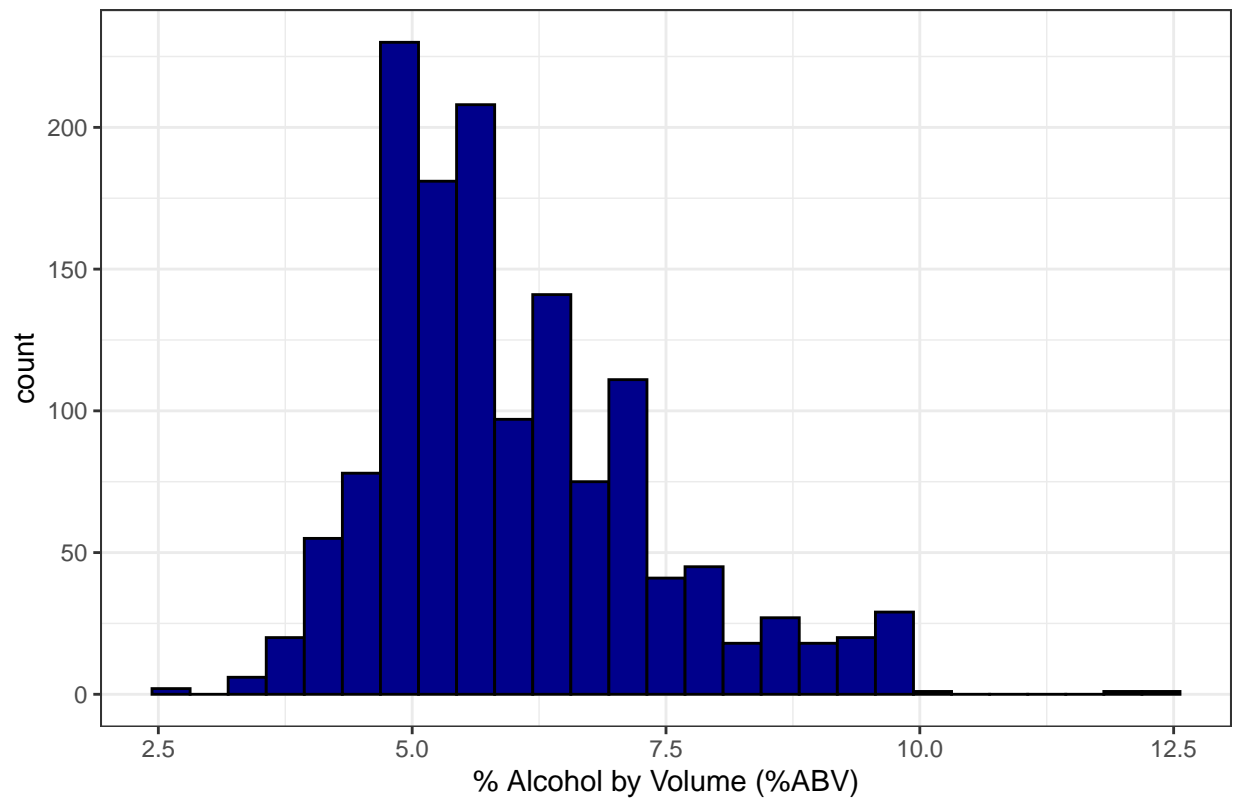
```
summary(cleanData$ABV)
```

```
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.02700 0.05000 0.05700 0.05991 0.06800 0.12500
```

```
# Histogram of ABV Percentage
```


```
cleanData %>% ggplot(aes(ABV*100)) + geom_histogram(fill="darkblue",color="black", binwidth= 0.375) +
```

Distribution of Beer %ABV, Right-Skewed




```
# Box Plot  
boxplot(cleanData$ABV, col='orange', main = 'Alcohol by volume')
```


Alcohol by volume



```
# QQ plot for normality check
qqnorm(cleanData$ABV, pch = 1, frame = FALSE)
qqline(cleanData$ABV, col = "steelblue", lwd = 2, main = 'Alcohol by volume')
```




We see it's quite right skewed from its histogram. QQ plot also showed that this is not normally distributed data as it has some curve at upper quantiles. In the box plot, we clearly see there are some outliers. such as London Balling beer we found from MAX ABV.

Relationship between ABV and IBU


We made scatter plot between ABV and IBU to see what is the relationship, and we see that it has some positive relationship that as ABV Value goes up IBU tend to go up as well.

```
# Scatter Plot for ABV vs IBU for each State
cleanData %>% ggplot(aes(x=IBU, y=ABV, color=State)) + geom_point()
```



```
#7 Scatter Plot for ABV vs IBU relationship
theme_set(theme_bw()) # pre-set the bw theme.
g <- ggplot(cleanData, aes(IBU, ABV, color='red'))
g + geom_point(color='blue') +
  geom_smooth(method="lm", se=F) +
  labs(y="ABV",
       x="IBU",
       title="ABV vs IBU Relationship")
```

```
## 'geom_smooth()' using formula 'y ~ x'
```



KNN Cluster Plot for American Ales

From the KNN cluster plot of Ale vs IPA, we found beers with high ABV and IBU are most likely IPA and beers with low ABV and IBU are most likely Ale. We want to know why there are some Ales on upper side of plot and IPAs on low IBU/ABV. To get in deeper, we made a KNN cluster plot just for Ale, especially American to reduce number of variables.

```
set.seed(4)
splitPerc = .70
aleData = cleanData %>% filter(grepl("Ale", Style) & grepl("American", Style))
trainIndices = sample(1:dim(aleData)[1], round(splitPerc * dim(aleData)[1]))
trainAle = aleData[trainIndices,]
testAle = aleData[-trainIndices,]
fit = knn(trainAle[,c(4,5)], testAle[,c(4,5)], trainAle$Style, k=6)

predAleDF = data.frame(testAle, predicted = fit)

predAleBoundary = data.frame(x = predAleDF$ABV,
                             y = predAleDF$IBU,
                             predicted = predAleDF$predicted)


find_hull = function(df) df[chull(df$x, df$y), ]
boundary = ddply(predAleBoundary, .variables = "predicted", .fun = find_hull)
```

```

palettes = brewer.pal(n = 9, name = "Set3")
colors = c("American Amber / Red Ale" = palettes[1], "American Black Ale" = palettes[2], "American Blonde Ale" = palettes[3], "American Brown Ale" = palettes[4], "American Dark Wheat Ale" = palettes[5], "American Pale Ale (APA)" = palettes[6], "American Pale Wheat Ale" = palettes[7], "American Strong Ale" = palettes[8], "American Wild Ale" = palettes[9])

ggplot() +
  geom_point(data=predAleDF, aes(ABV, IBU, color=predicted, fill=predicted), size = 5) +
  geom_polygon(data = boundary, aes(x,y, color=predicted, fill=predicted), alpha = 0.5) +
  geom_point(aes(ABV, IBU, color=Style), data=testAle) + ggtitle("KNN Cluster Plot for American Ales")

```



Looking at the plot, we see that some types of Ale have high IBU and ABV even though it's not an IPA. American black Ale is one of the types, which is sometimes called as Black IPA. If we set this type as IPA, we could have gotten better classification results than before.

KNN Cluster Plot for IPA

We made another KNN cluster plot just for IPA.

```

set.seed(4)
splitPerc = .70
ipaData = cleanData %>% filter(grepl("IPA", Style))
trainIndices = sample(1:dim(ipaData)[1], round(splitPerc * dim(ipaData)[1]))
trainIPA = ipaData[trainIndices,]
testIPA = ipaData[-trainIndices,]
fit = knn(trainIPA[,c(4,5)], testIPA[,c(4,5)], trainIPA$Style, k=6)
predIpaDF = data.frame(testIPA, predicted = fit)

```


```

predIpaBoundary = data.frame(x = predIpaDF$ABV,
                             y = predIpaDF$IBU,
                             predicted = predIpaDF$predicted)

find_hull = function(df) df[chull(df$x, df$y), ]
boundary = ddply(predIpaBoundary, .variables = "predicted", .fun = find_hull)

colors = c("American Double / Imperial IPA" = "red", "American IPA" = "blue", "Belgian IPA" = "green", "English India Pale Ale (IPA)" = "yellow")
ggplot() +
  geom_point(data=predIpaDF, aes(ABV, IBU, color=predicted, fill=predicted), size = 5) +
  geom_polygon(data = boundary, aes(x,y, color=predicted, fill=predicted), alpha = 0.5) +
  geom_point(aes(ABV, IBU, color=Style), data=testIPA) + ggtitle("KNN Cluster Plot") + ggtitle("KNN Cluster Plot")

```



We found American IPA have much broader range of IBU and ABV that they contains. With just two information, IBU and ABV, it is not enough to understand the relationship of ABV and IBU against its style. For IPA, we may need additional feature variables such as hop ratio or type of ingredients to get better classification model.

From the above chart, we found that Colarado and Orgeon have more high ABV and IBU Ales than other states. Thus we can make a question that why these states have more Ales that have high bitterness and alcohol level.